

Estimating the Mean from Data: Introduction to Estimation Theory

Motonobu Kanagawa

Introduction to Statistics, EURECOM

March 4, 2024

Outline

Mean Estimation Problem and Motivations

The Data Generation Process Matters

Preliminaries: Key Properties of Expectation and Variance

Statistical Estimators

Mean Square Error and Bias-Variance Decomposition

Bias-Variance Decomposition in Mean Estimation

Consistency and Unbiasedness

Variance Reduction by Introducing a Bias

Estimation of the Mean

- Let X be a random variable taking values in \mathbb{R} with probability distribution P .

(Note: In the previous lecture, P is used to denote the distribution of the underlying probability space, but here P denotes the distribution of X).

- The **mean** (or the **expected value**) of X is defined by

$$\mu := \mathbb{E}_{X \sim P}[X] = \int x \, dP(x) \in \mathbb{R}.$$

Assume that **we don't know P** , and thus **we don't know μ** .

Estimation of the Mean

- Assume instead that we are given some **data**:

$$X_1, \dots, X_n \in \mathbb{R}$$

- These are assumed to be **random variables** taking values in \mathbb{R} .
- The task of mean estimation is **estimating the unknown mean μ** from **the data X_1, \dots, X_n** .
- This is one of the most ubiquitous and fundamental problems in statistics.
- In this lecture, we look at this problem in details.

Motivation 1: Relation to Many Problems

Many problems can be formulated as estimation of the mean.

Examples:

- Monte Carlo: Simulation-based mean estimation.
- Design of experiments: Average treatment (causal) effect.
- Regression: Estimation of the conditional mean.
- Supervised machine learning:
 - ▶ Risk = the mean of a loss function.
 - ▶ Stochastic gradient = approximation of the expected gradient.

Motivation 2: Different Statistical Approaches

Mean estimation can be used for illustrating different approaches.

- The “frequentist” approach - maximum likelihood estimation.
- The “Bayesian” approach - posterior inference.
- The “empirical Bayes” - the mixed approach.

Motivation 3: Key Notions

We can learn **key notions** in statistics.

- Estimator and consistency.
- Bias-variance decomposition/trade-off
- Law of large numbers and the central limit theorem.

Most importantly,

- The key is **how data are generated/obtained**.

Is the Empirical Average a Good Approach?

A standard approach is to take the **empirical average** of data points:

$$\hat{\mu} := \frac{1}{n} \sum_{i=1}^n X_i.$$

In this lecture, we will address questions like:

- When is the empirical average a **good estimate**, and when is it not?
- When can we **justify** the use of the empirical average?
- What **conditions** do we need for the data X_1, \dots, X_n ?

Outline

Mean Estimation Problem and Motivations

The Data Generation Process Matters

Preliminaries: Key Properties of Expectation and Variance

Statistical Estimators

Mean Square Error and Bias-Variance Decomposition

Bias-Variance Decomposition in Mean Estimation

Consistency and Unbiasedness

Variance Reduction by Introducing a Bias

Population and Data

In the mean estimation problem, we have two kinds of **random variables**:

[Population] Random variable X represents the **hypothetical population** of interest, with P being its probability distribution.

[Data] Random Variables X_1, \dots, X_n represent the **given data**.

- The data X_1, \dots, X_n are assumed to provide **information** about the population random variable X (or its distribution P).
- Otherwise, we cannot estimate the population mean $\mu = \mathbb{E}_{X \sim P}[X]$ from the data X_1, \dots, X_n .
- Therefore, **how the data are generated/obtained** becomes very important.

Example: Estimating the Average Income in France

- Assume that $X \in \mathbb{R}$ represents the income of a **randomly sampled** French person, with P being its distribution.
- The population mean $\mu = \mathbb{E}_{X \sim P}[X]$ represents the average income of French people.
- The data X_1, \dots, X_n are the incomes of n French people **randomly selected** from the French population.
- Then, is the empirical average

$$\hat{\mu} := \frac{1}{n} \sum_{i=1}^n X_i$$

a good estimate of the true average income $\mu = \mathbb{E}_{X \sim P}[X]$?

Example: Estimating the Average Income in France

- Assume that data X_1, \dots, X_n are the incomes of randomly sampled French persons in French Riviera.
- Then, the empirical average

$$\hat{\mu} := \frac{1}{n} \sum_{i=1}^n X_i$$

would be higher than the average income of the French population.

Example: Estimating the Average Income in France

- Assume that the data X_1, \dots, X_n are the incomes of randomly sampled French people **between age 20 and 30**.
- Then, the empirical average

$$\hat{\mu} := \frac{1}{n} \sum_{i=1}^n X_i$$

would give an estimate **lower** than the true average income.

The Data Generating Process Matters

These examples indicate that **how the data are generated/obtained strongly affects the validity** of the empirical average.

- We need to make sure that data X_1, \dots, X_n are sampled from the **same population** as that of the target random variable $X \sim P$.
- This requirement is mathematically formulated by assuming that random variables X_1, \dots, X_n are **independently and identically distributed (i.i.d.)** with $X \sim P$.

Independently and Identically Distributed (i.i.d.)

Recall that random variables X_1, \dots, X_n are i.i.d. with a random variable $X \sim P$ if they satisfy the following:

- Independence:

- ▶ X_i and X_j are **independent** for all $i \neq j$.
- ▶ X_i and X are **independent** for all $i = 1, \dots, n$;
 - ▶ Recall that X represents the hypothetical population (e.g., randomly selected French person).

- Identity:

- ▶ X_i follows the **same probability distribution** P of X (for all $i = 1, \dots, n$).

We often write $X_1, \dots, X_n \sim P$ (i.i.d.).

See also the lecture slides on Probability Theory.

Outline

Mean Estimation Problem and Motivations

The Data Generation Process Matters

Preliminaries: Key Properties of Expectation and Variance

Statistical Estimators

Mean Square Error and Bias-Variance Decomposition

Bias-Variance Decomposition in Mean Estimation

Consistency and Unbiasedness

Variance Reduction by Introducing a Bias

Preliminaries

Before going further, we collect here some key properties of **Expectation** and **Variance** of random variables.

Some Key Properties of Expectation

- For any real-valued random variable X and a constant $c \in \mathbb{R}$, we have

$$\mathbb{E}[cX] = c\mathbb{E}[X].$$

- For any real-valued random variables X and Y , we have

$$\mathbb{E}[X + Y] = \mathbb{E}[X] + \mathbb{E}[Y].$$

- If X and Y are independent,

$$\mathbb{E}[XY] = \mathbb{E}[X]\mathbb{E}[Y]$$

Variance of a Random Variable

In statistics, the **variance** of a random variable plays a key role.

- Let X be a real-valued random variable with probability distribution P .

Then the variance of X is defined by

$$\mathbb{V}[X] := \mathbb{E}[(X - \mathbb{E}[X])^2] = \int (x - \mathbb{E}[X])^2 dP(x) \geq 0.$$

- Note that the mean $\mathbb{E}[X] \in \mathbb{R}$ is a constant.

Some Key Properties of Variance

Let X be a real-valued random variable.

Then we have

$$\mathbb{V}[X] = \mathbb{E}[X^2] - (\mathbb{E}[X])^2$$

Proof:

$$\begin{aligned}\mathbb{V}[X] &= \mathbb{E}[(X - \mathbb{E}[X])^2] \\ &= \mathbb{E}[X^2 - 2X\mathbb{E}[X] + (\mathbb{E}[X])^2] \\ &= \mathbb{E}[X^2] - 2(\mathbb{E}[X])^2 + (\mathbb{E}[X])^2 \\ &= \mathbb{E}[X^2] - (\mathbb{E}[X])^2\end{aligned}$$

Some Key Properties of Variance

Let X be a real-valued random variable. Then for any constant $c \in \mathbb{R}$, we have

$$\mathbb{V}[cX] = c^2 \mathbb{V}[X].$$

Proof:

$$\begin{aligned}\mathbb{V}[cX] &:= \mathbb{E}[(cX - \mathbb{E}[cX])^2] \\ &= c^2 \mathbb{E}[(X - \mathbb{E}[X])^2] = c^2 \mathbb{V}[X].\end{aligned}$$

In particular, by setting $c = 1/n$, we have

$$\mathbb{V}\left[\frac{X}{n}\right] = \frac{1}{n^2} \mathbb{V}[X].$$

Some Key Properties of Variance

Let X and Y be real-valued random variables.

If X and Y are **independent**, then

$$\mathbb{V}[X + Y] = \mathbb{V}[X] + \mathbb{V}[Y].$$

Proof:

$$\begin{aligned}\mathbb{V}[X + Y] &:= \mathbb{E}[(X + Y - \mathbb{E}[X + Y])^2] \\&= \mathbb{E}[(X - \mathbb{E}[X] + Y - \mathbb{E}[Y])^2] \\&= \mathbb{E}[(X - \mathbb{E}[X])^2 + 2(X - \mathbb{E}[X])(Y - \mathbb{E}[Y]) + (Y - \mathbb{E}[Y])^2] \\&= \mathbb{E}[(X - \mathbb{E}[X])^2] + 2\mathbb{E}[(X - \mathbb{E}[X])(Y - \mathbb{E}[Y])] + \mathbb{E}[(Y - \mathbb{E}[Y])^2] \\&= \mathbb{E}[(X - \mathbb{E}[X])^2] + \mathbb{E}[(Y - \mathbb{E}[Y])^2] = \mathbb{V}[X] + \mathbb{V}[Y],\end{aligned}$$

where we used

$$\mathbb{E}[(X - \mathbb{E}[X])(Y - \mathbb{E}[Y])] = \mathbb{E}[(X - \mathbb{E}[X])]\mathbb{E}[(Y - \mathbb{E}[Y])] = 0,$$

which follows from the **independence** of X and Y .

Some Key Properties of Variance

By recursive applications of the previous result, we have the following useful result:

Let X_1, X_2, \dots, X_n are **independent** real-valued random variables (note: they don't necessary **identically** distributed).

Then we have

$$\mathbb{V}[\sum_{i=1}^n X_i] = \sum_{i=1}^n \mathbb{V}[X_i]$$

Corollary:

- ▶ Let X_1, \dots, X_n be independent real-valued random variables.
- ▶ Let $c_1, \dots, c_n \in \mathbb{R}$ be constants.

Then
*Weighted
average*

$$\mathbb{V}[\sum_{i=1}^n c_i X_i] = \sum_{i=1}^n \mathbb{V}[c_i X_i] = \sum_{i=1}^n c_i^2 \mathbb{V}[X_i].$$

Some Key Properties of Variance

In particular, assuming that X_1, \dots, X_n are **i.i.d.** with a random variable X , and setting $c_i := 1/n$, we have

$$\mathbb{V}\left[\frac{1}{n} \sum_{i=1}^n X_i\right] = \frac{1}{n^2} \sum_{i=1}^n \mathbb{V}[X_i] = \frac{1}{n} \mathbb{V}[X].$$

*n times the
same
variance
because iid*

- Thus, the variance of the empirical average $\frac{1}{n} \sum_{i=1}^n X_i$ is n times smaller than the variance of X .
- By taking the average over **independent** observations, the variance can be reduced.

Outline

Mean Estimation Problem and Motivations

The Data Generation Process Matters

Preliminaries: Key Properties of Expectation and Variance

Statistical Estimators

Mean Square Error and Bias-Variance Decomposition

Bias-Variance Decomposition in Mean Estimation

Consistency and Unbiasedness

Variance Reduction by Introducing a Bias

Estimators and Estimates

In statistics, the procedure of estimating a quantity of interest is formulated as a function of data.

- This function is called an estimator.
- The output from the estimator is called an estimate.

Estimators and Estimates

- Let $\theta^* \in \Theta$ be an **unknown quantity of interest** that we want to estimate (Θ is an appropriate set) (θ^* is also called an **estimand**).
- Assume that we are given some **data** D_n of size $n \in \mathbb{N}$ of the form

$$D_n := (X_1, \dots, X_n) \in \mathcal{X}^n \quad (\text{X times, X times, } \dots)$$

where each $X_i \in \mathcal{X}$ is a random variable (\mathcal{X} is a measurable space.).

Definition: a map

$$F_n : \mathcal{X}^n \rightarrow \Theta$$

is called an **estimator** (of θ^*).

- The estimator should be designed so that the estimate will be close to θ^* .
- $\hat{\theta}_n := F_n(D_n)$ is called an **estimate** (of θ^*).

Estimators and Estimates: Mean Estimation

Let's consider the **mean estimation problem** as an example.

The quantity of interest is the **mean** of the random variable $X \sim P$:

$$\theta^* := \mu := \mathbb{E}[X] \in \mathbb{R} =: \Theta.$$

Assume that n random variables X_1, \dots, X_n are given as **data**:

$$D_n = (X_1, \dots, X_n) \in \mathcal{X}^n, \quad \mathcal{X} := \mathbb{R}.$$

Then one can define an estimator $F_n : \mathcal{X}^n \rightarrow \Theta$ of the mean θ^* by

$$F_n(D_n) := \frac{1}{n} \sum_{i=1}^n X_i = \hat{\mu} =: \hat{\theta}.$$

i.e., the empirical average of X_1, \dots, X_n .

Which Estimator Should We Choose?

Note that the empirical average is not the only choice.

For instance, we can define various estimators for the mean estimation problem; e.g.,

1. $F_n(D_n) := (X_1 + \dots + X_n)/n.$

empirical average

2. $F_n(D_n) := X_1$ (i.e., discarding X_2, \dots, X_n).

3. $F_n(D_n) := 0$ (i.e., always outputs constant 0, no matter what D_n is).

4. $F_n(D_n) := c_0 + c_1 X_1 + \dots + c_n X_n$ for some $c_0, c_1, \dots, c_n \geq 0$.

or define weighted average

- Which estimator should we choose?
- When is the empirical average a good choice, and when is it not?

(Actually we'll see that the empirical average is **not always** a good choice).

Which Estimator Should We Choose?

To investigate these questions, we need to introduce **criteria** for comparing different estimators.

Outline

Mean Estimation Problem and Motivations

The Data Generation Process Matters

Preliminaries: Key Properties of Expectation and Variance

Statistical Estimators

Mean Square Error and Bias-Variance Decomposition

Bias-Variance Decomposition in Mean Estimation

Consistency and Unbiasedness

Variance Reduction by Introducing a Bias

Mean Square Error (MSE)

To discuss the quality of a statistical estimator, we need a certain **error criterion**.

Here we consider the **mean square error (MSE)**, one of the most standard criteria.

- Let $\theta^* \in \Theta \subset \mathbb{R}$ be the unknown quantity of interest.
- We assume $\Theta \subset \mathbb{R}$ for simplicity, but the following argument also holds for more general situations.
- Consider an estimator $F_n : \mathcal{X}^n \rightarrow \Theta$ such that

$$\hat{\theta}_n := F_n(D_n) \in \Theta, \quad D_n := (X_1, \dots, X_n) \in \mathcal{X}^n.$$

- Note that the estimate $\hat{\theta}_n = F_n(D_n) = F_n((X_1, \dots, X_n))$ is a **random variable**, since X_1, \dots, X_n are random variables.

Mean Square Error (MSE)

- Then we can consider the **squared error** between the target θ^* and estimate $\hat{\theta}_n$:

$$(\hat{\theta}_n - \theta^*)^2 = (F_n(D_n) - \theta^*)^2.$$

- This error is also a random variable, because the estimate $\hat{\theta}_n = F_n(D_n)$ is a random variable.
- Then the **mean square error (MSE)** of the estimator F_n is defined as the expectation of the squared error:

$$\mathbb{E}[(\hat{\theta}_n - \theta^*)^2] = \mathbb{E}[(F_n(D_n) - \theta^*)^2]$$

where the expectation is with respect to the **data** $D_n = (X_1, \dots, X_n)$.

- The MSE quantifies how the estimate $\hat{\theta}_n$ is close to (or far from) the target θ^* **on average**.

Mean Square Error (MSE)

Note that the MSE

$$\mathbb{E}[(\hat{\theta}_n - \theta^*)^2] = \mathbb{E}[(F_n(D_n) - \theta^*)^2]$$

depends on

1. the target quantity θ^*
2. the estimator F_n
3. the **distribution of the data** X_1, \dots, X_n

By theoretically studying the MSE, we can study

- ▶ which estimator F_n is good for estimating the target θ^* ,
- ▶ when the data X_1, \dots, X_n are **distributed in an assumed way**.

Probabilistic Error Bound from MSE

- A general fact: For any **non-negative** real-valued random variable Z , **Markov's inequality** states that

$$\Pr(Z \geq c) \leq \frac{\mathbb{E}[Z]}{c}, \quad \forall c > 0.$$

- By setting $Z := (\hat{\theta}_n - \theta^*)^2$, we then have

$$\Pr((\hat{\theta}_n - \theta^*)^2 \geq c) \leq \frac{\mathbb{E}[(\hat{\theta}_n - \theta^*)^2]}{c}, \quad \forall c > 0.$$

- Thus, if the MSE $\mathbb{E}[(\hat{\theta} - \theta^*)^2]$ is small, then the probability of

$$(\hat{\theta}_n - \theta^*)^2 > c$$

becomes small for any $c > 0$.

Bias-Variance Decomposition

- The following is a **very important** result concerning the MSE.

Theorem: The MSE can be decomposed into the **bias** and the **variance** of the estimator, as follows:

$$\mathbb{E}[(\hat{\theta}_n - \theta^*)^2] = \underbrace{\mathbb{E}[(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])^2]}_{\text{Variance}} + \underbrace{(\mathbb{E}[\hat{\theta}_n] - \theta^*)^2}_{\text{Bias}}$$

This is called the **bias-variance decomposition**.

- The **bias** of the estimator $F_n : \mathcal{X}^n \rightarrow \Theta$ is defined as the **difference** between the **expectation of the estimate** $\mathbb{E}[\hat{\theta}_n]$ and the **target** θ^* :

$$\mathbb{E}[\hat{\theta}_n] - \theta^* = \mathbb{E}[F_n(D_n)] - \theta^*.$$

where the expectation is with respect to the data $D_n = (X_1, \dots, X_n)$.

Bias-Variance Decomposition

$$\mathbb{E}[(\hat{\theta}_n - \theta^*)^2] = \underbrace{\mathbb{E}[(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])^2]}_{\text{Variance}} + \underbrace{(\mathbb{E}[\hat{\theta}_n] - \theta^*)^2}_{\text{Bias}}$$

- The **variance** of the estimator $F_n : \mathcal{X}^n \rightarrow \Theta$ is defined as

$$\mathbb{E}[(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])^2] = \mathbb{E}[(F_n(D_n) - \mathbb{E}[F_n(D_n)])^2].$$

- ▶ i.e., the **average deviation** of the estimate $\hat{\theta}_n := F_n(D_n)$ from its mean $\mathbb{E}[\hat{\theta}_n]$.
 - ▶ Recall again that the estimate $\hat{\theta}_n$ is a random variable.
- To make the mean-square error small, **both the bias and variance need to be small!**

Proof of Bias-Variance Decomposition

- The mean square error can be expanded as

$$\mathbb{E}[(\hat{\theta}_n - \theta^*)^2]$$

$$= \mathbb{E}[(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n] + \mathbb{E}[\hat{\theta}_n] - \theta^*)^2]$$

$$= \mathbb{E}[(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])^2] + \mathbb{E}[(\mathbb{E}[\hat{\theta}_n] - \theta^*)^2] + 2\mathbb{E}[(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])(\mathbb{E}[\hat{\theta}_n] - \theta^*)]$$

$$= \mathbb{E}[(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])^2] + (\mathbb{E}[\hat{\theta}_n] - \theta^*)^2,$$

where the last line follows from $\mathbb{E}[\hat{\theta}_n]$ being a constant:

$$\mathbb{E}[(\mathbb{E}[\hat{\theta}_n] - \theta^*)^2] = (\mathbb{E}[\hat{\theta}_n] - \theta^*)^2,$$

$$\mathbb{E}[(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])(\mathbb{E}[\hat{\theta}_n] - \theta^*)] = \mathbb{E}[(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])] (\mathbb{E}[\hat{\theta}_n] - \theta^*) = 0.$$

Constant can be extracted

Remarks on the Bias-Variance Decomposition

- The bias-variance decomposition holds under a very generic situation.
 - ▶ This is because the proof does **not** require any assumption about the **joint distribution of the data** X_1, \dots, X_n (essentially).
 - ▶ The only assumption is that the MSE is finite.
- Thus, for instance, we can consider cases like:
 - ▶ where X_1, \dots, X_n are **not independently** distributed
 - ▶ where X_1, \dots, X_n are **not identically** distributed.
- By considering a different setting for the distribution of the data X_1, \dots, X_n , we can study **when** a certain estimator is a good choice, **when** it is not.
- This is done by analyzing the bias and variance of the estimator.

Bias-Variance Decomposition: Multivariate Case

- Let $\theta^* \in \Theta \subset \mathbb{R}^d$ be the quantity of interest.
- Let $\hat{\theta}_n$ be any estimate of θ^* (you can just think of $\hat{\theta}_n$ as a random variable in \mathbb{R}^d).
- Define the mean square error by

$$\mathbb{E}[\|\hat{\theta}_n - \theta^*\|^2],$$

where $\|\cdot\|$ is the norm of \mathbb{R}^d .

Theorem. - Assume that

$$\|\mathbb{E}[\hat{\theta}_n]\| < \infty, \quad \mathbb{E}[\|\hat{\theta}_n\|^2] < \infty.$$

Then the following bias-variance decomposition holds:

$$\mathbb{E}[\|\hat{\theta}_n - \theta^*\|^2] = \underbrace{\mathbb{E}[\|\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n]\|^2]}_{\text{Variance}} + \underbrace{\|\theta^* - \mathbb{E}[\hat{\theta}_n]\|^2}_{\text{Bias}}$$

Bias-Variance Decomposition: Multivariate Case

Exercise: Prove the above bias-variance decomposition.

Hint: for any $a, b \in \mathbb{R}^d$,

$$\|a - b\|^2 = \langle a - b, a - b \rangle$$

where $\langle \cdot, \cdot \rangle$ is the inner product in \mathbb{R}^d .

Outline

Mean Estimation Problem and Motivations

The Data Generation Process Matters

Preliminaries: Key Properties of Expectation and Variance

Statistical Estimators

Mean Square Error and Bias-Variance Decomposition

Bias-Variance Decomposition in Mean Estimation

Consistency and Unbiasedness

Variance Reduction by Introducing a Bias

Mean Estimation Problem: Setup

- Now consider the **mean estimation** problem.
- Let $X \sim P$ be the random variable of interest, whose mean

$$\mu_P := \mathbb{E}[X] = \int x \, dP(x)$$

is the estimand.

- To deal with a generic situation, we assume that i.i.d. data X_1, \dots, X_n are generated from a **probability distribution** Q , which **can be different from** P :

$$X_1, \dots, X_n \sim Q, i.i.d.$$

- Let $Y \sim Q$ be a random variable, with distribution Q ;
- Then X_1, \dots, X_n are i.i.d. with Y .

Bias-Variance Decomposition in Mean Estimation

- Assume that the mean and the variance of $Y \sim Q$ are finite:

$$|\mu_Q| < \infty, \quad \mu_Q := \mathbb{E}_{Y \sim Q}[Y]$$
$$\sigma_Q^2 < \infty, \quad \sigma_Q^2 := \mathbb{V}_{Y \sim Q}[Y] := \mathbb{E}_{Y \sim Q}[(Y - \mu_Q)^2].$$

Theorem: The mean square error of the empirical average estimator

$$\hat{\mu} := \frac{1}{n} \sum_{i=1}^n X_i$$

is given by

$$\begin{aligned} \mathbb{E}[(\hat{\mu} - \mu_P)^2] &= \mathbb{E}[(\hat{\mu} - \mathbb{E}[\hat{\mu}])^2] + (\mathbb{E}[\hat{\mu}] - \mu_P)^2 \\ &= \frac{\sigma_Q^2}{n} + (\mu_Q - \mu_P)^2. \end{aligned}$$

Proof: Bias-Variance Decomposition in Mean Estimation

Proof:

- The first identity follows from the bias-variance decomposition.
- Thus, we show the second identity.

Variance term.

Because X_1, \dots, X_n are i.i.d. with $Y \sim Q$, the variance term can be expressed as

$$\begin{aligned}\mathbb{E}[(\hat{\mu} - \mathbb{E}[\hat{\mu}])^2] &= \mathbb{V}[\hat{\mu}] = \mathbb{V}\left[\frac{1}{n} \sum_{i=1}^n X_i\right] \\ &= \sum_{i=1}^n \mathbb{V}\left[\frac{1}{n} X_i\right] = \frac{1}{n^2} \sum_{i=1}^n \mathbb{V}[X_i] = \frac{1}{n} \mathbb{V}[Y] = \frac{\sigma_Q^2}{n}.\end{aligned}$$

Proof: Bias-Variance Decomposition in Mean Estimation

Bias term. On the other hand,

$$\mathbb{E}[\hat{\mu}] = \mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n X_i\right] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}[X_i] = \mathbb{E}[Y] = \mu_Q.$$

Therefore, the bias term is

$$(\mathbb{E}[\hat{\mu}] - \mu_P)^2 = (\mu_Q - \mu_P)^2.$$



Interpretation of the Bias-Variance Decomposition

We proved the bias-variance decomposition:

$$\mathbb{E}[(\hat{\mu} - \mu_P)^2] = \frac{\sigma_Q^2}{n} + (\mu_Q - \mu_P)^2.$$

Let's study what this means.

- The bias of the estimator $\hat{\mu} = \frac{1}{n} \sum_{i=1}^n X_i$ is

$$\mathbb{E}[\hat{\mu}] - \mu_P = \mu_Q - \mu_P$$

i.e., the difference between

- ▶ the mean μ_Q of the data distribution Q , and
- ▶ the mean μ_P of the target distribution P .

Interpretation of the Bias-Variance Decomposition

Therefore,

- ▶ if the data X_1, \dots, X_n are independently generated from a distribution Q , and
- ▶ if the mean μ_Q of Q is **different** from the mean μ_P of the target random variable $X \sim P$,

then the use of the empirical average

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n X_i,$$

causes a **non-zero bias**, $\mu_Q - \mu_P \neq 0$.

Note that in this case, since $(\mu_Q - \mu_P)^2 > 0$, the mean square error

$$\mathbb{E}[(\hat{\mu} - \mu_P)^2] = (\mu_Q - \mu_P)^2 + \frac{\sigma_Q^2}{n} \geq (\mu_Q - \mu_P)^2 > 0$$

does **not decrease to 0**, even when $n \rightarrow \infty$.

Interpretation of the Bias-Variance Decomposition

This example shows the importance of the **data distribution** Q .

- If possible, we should collect data X_1, \dots, X_n generated from the **same distribution** P as the target random variable X , i.e., $Q = P$.
- In this case, the bias becomes 0: $(\mu_Q - \mu_P)^2 = 0$, and the MSE is

$$\mathbb{E}[(\hat{\mu} - \mu_P)^2] = \frac{\sigma_P^2}{n},$$

where $\sigma_P^2 = \mathbb{V}[X]$ is the variance of $X \sim P$.

- Thus, the MSE decreases as the sample size n increases.

Interpretation of the Bias-Variance Decomposition

- On the other hand, the variance term

$$\mathbb{E}[(\hat{\mu} - \mathbb{E}[\hat{\mu}])^2] = \frac{\sigma_Q^2}{n}$$

depends only on the data X_1, \dots, X_n , and not on the target μ_P .

- Therefore, whatever the data distribution Q is, the variance term converges to 0 as $n \rightarrow \infty$:

$$\lim_{n \rightarrow \infty} \mathbb{E}[(\hat{\mu} - \mathbb{E}[\hat{\mu}])^2] = \lim_{n \rightarrow \infty} \frac{\sigma_Q^2}{n} = 0.$$

Interpretation of the Bias-Variance Decomposition

- Note that in the derivation of the variance term, we used

$$\mathbb{V}[\frac{1}{n} \sum_{i=1}^n X_i] = \sum_{i=1}^n \mathbb{V}[\frac{1}{n} X_i].$$

- This follows from the **independence** between X_1, \dots, X_n . (see pp.21-22)
- Therefore, if the independence between X_1, \dots, X_n does **not** hold, the variance **may not decrease to 0** (we'll see an example later).

Interpretation of the Bias-Variance Decomposition

- For example, recall the example where $X \sim P$ represents the income of a randomly picked-up French person.
- Assume that data $X_1, \dots, X_n \sim Q$ (*i.i.d.*) are the incomes of randomly picked-up French persons in French Riviera.
- Then we would have

$$\mu_Q := \mathbb{E}_{Y \sim Q}[Y] > \mathbb{E}_{X \sim P}[X] =: \mu_P$$

i.e., the average income of French Riviera people μ_Q is higher than the average income of the whole population μ_P .

Interpretation of the Bias-Variance Decomposition

- Thus, the empirical average of the data

$$\hat{\mu} := \frac{1}{n} \sum_{i=1}^n X_i$$

has a **non-zero bias**:

$$\mathbb{E}[\hat{\mu}] - \mu_P = \mu_Q - \mu_P \neq 0.$$

- Therefore, the MSE of the empirical average

$$\mathbb{E}[(\hat{\mu} - \mu_P)^2] = (\mu_Q - \mu_P)^2 + \frac{\sigma_Q^2}{n}$$

does **not decrease to 0**, even when n is **very large**.

- Thus, we should make sure that data X_1, \dots, X_n are randomly picked-up from the **whole French population**. (i.e., $Q = P$).

Mean Estimation in the Multivariate Case

- Let $X \sim P$ be a random vector in \mathbb{R}^d . Define

$$\mu_P := \mathbb{E}_{X \sim P}[X] \in \mathbb{R}^d$$

- Let $X_1, \dots, X_n \sim Q$ (*i.i.d.*) be random vectors in \mathbb{R}^d , and let $Y \sim Q$. Define

$$\mu_Q := \mathbb{E}_{Y \sim Q}[Y] \in \mathbb{R}^d, \quad \sigma_Q^2 := \mathbb{E}_{Y \sim Q}[\|Y - \mu_Q\|^2] \geq 0.$$

Theorem. Assume that

$$\|\mu_P\| < \infty, \quad \|\mu_Q\| < \infty, \quad \sigma_Q^2 < \infty.$$

Then, the empirical average estimator $\hat{\mu} = \frac{1}{n} \sum_{i=1}^n X_i$ satisfies

$$\begin{aligned} \mathbb{E}[\|\hat{\mu} - \mu_P\|^2] &= \mathbb{E}[\|\hat{\mu} - \mathbb{E}[\hat{\mu}]\|^2] + \|\mathbb{E}[\hat{\mu}] - \mu_P\|^2 \\ &= \frac{\sigma_Q^2}{n} + \|\mu_Q - \mu_P\|^2. \end{aligned}$$

Exercise. Prove this. (The first identity is the bias-variance decomposition)

How Large should the Sample Size be?

- In the mean estimation problem, when $X_1, \dots, X_n \sim P$ i.i.d., the MSE is given by

$$\mathbb{E}[(\hat{\mu} - \mu_P)^2] = \frac{\sigma_P^2}{n}, \quad \sigma_P^2 := \mathbb{V}[X].$$

for the empirical average estimate $\hat{\mu} := \frac{1}{n} \sum_{i=1}^n X_i$.

- Assume that one wants to make the MSE small in that sense that

$$\mathbb{E}[(\hat{\mu} - \mu_P)^2] \leq \varepsilon^2,$$

for some $\varepsilon > 0$. Then the sample size n should satisfy

$$\mathbb{E}[(\hat{\mu} - \mu_P)^2] = \frac{\sigma_P^2}{n} \leq \varepsilon^2$$

or equivalently

$$n \geq \frac{\sigma_P^2}{\varepsilon^2}.$$

How Large should the Sample Size be?

For instance, consider the example of estimating the average income.

- Assume $\mu_P = 2,000$ EUR/month (mean) and $\sigma_P = 500$ (standard deviation).

Then, the sample size n should satisfy

$$n \geq \frac{500^2}{\varepsilon^2}.$$

For instance,

- to achieve the precision of $\varepsilon = 10$, we need $n \geq 2500$.
- to achieve the precision of $\varepsilon = 1$, we need $n \geq 250,000$.

Outline

Mean Estimation Problem and Motivations

The Data Generation Process Matters

Preliminaries: Key Properties of Expectation and Variance

Statistical Estimators

Mean Square Error and Bias-Variance Decomposition

Bias-Variance Decomposition in Mean Estimation

Consistency and Unbiasedness

Variance Reduction by Introducing a Bias

Consistency

- Let $\theta^* \in \Theta \subset \mathbb{R}$ be an estimand (i.e., the quantity of interest).
- Let X_1, \dots, X_n be random variables such that $X_i \in \mathcal{X}$, and define the data as

$$D_n := (X_1, \dots, X_n) \in \mathcal{X}^n$$

- Let $F_n : \mathcal{X}^n \rightarrow \mathbb{R}$ be an estimator, and let $\hat{\theta}_n := F_n(D_n)$ be an estimate.

Definition. We call F_n a **consistent estimator** of θ^* , if the **estimate $\hat{\theta}_n$ converges to θ^* as $n \rightarrow \infty$** in an appropriate sense, e.g.,

$$\mathbb{E}[(\hat{\theta}_n - \theta^*)^2] \rightarrow 0 \quad \text{as } n \rightarrow \infty.$$

- The consistency means that, as we have **more data X_1, \dots, X_n** , the estimate $\hat{\theta}_n$ becomes **more accurate** (in estimating θ^*).
- Consistency is one of the **most important** concepts in statistics.

Unbiasedness

Definition. We call F_n an **unbiased estimator** of θ^* , if the bias is zero for every $n \in \mathbb{N}$, i.e.,

$$\mathbb{E}[F_n(D_n)] - \theta^* = \mathbb{E}[\hat{\theta}_n] - \theta^* = 0, \quad \forall n \in \mathbb{N}.$$

- If this is not satisfied, we call F_n a **biased** estimator of θ^* .

Unbiasedness

For instance, consider the **mean estimation** problem.

- If the data X_1, \dots, X_n are **i.i.d. with $X \sim P$** , then the empirical average $\hat{\mu} := \frac{1}{n} \sum_{i=1}^n X_i$ satisfies

$$\mathbb{E}[\hat{\mu}] = \mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n X_i\right] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}[X_i] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}[X] = \mathbb{E}[X] = \mu_P.$$

- So, in this case, the empirical average $\hat{\mu}$ is an **unbiased** estimator of the mean μ_P .
- If X_1, \dots, X_n are **i.i.d. with $Y \sim Q$** , and if $\mu_Q \neq \mu_P$, then

$$\mathbb{E}[\hat{\mu}] = \mu_Q \neq \mu_P.$$

- So, in this case, the empirical average $\hat{\mu}$ is a **biased** estimator of the mean μ_P .

Unbiasedness

- If F_n is an **unbiased** estimator, then the MSE is given by

$$\mathbb{E}[(\hat{\theta}_n - \theta^*)^2] = \mathbb{E}[(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])^2] = \mathbb{V}[\hat{\theta}_n]$$

i.e., the MSE is equal to the **variance** of the estimate $\hat{\theta}_n$.

Some important consequences of unbiasedness:

- If the **variance** $\mathbb{V}[\hat{\theta}_n]$ decreases to 0 as $n \rightarrow \infty$, then $\hat{\theta}_n$ converges to θ^* ; thus F_n becomes a **consistent** estimator.
- If we can **estimate** the variance $\mathbb{V}[\hat{\theta}_n]$, then we can **estimate the amount of error (MSE)**:
 - ▶ In other words, we can estimate **how far** the estimate $\hat{\theta}_n$ is from the target θ^* .
 - ▶ Thus, an **estimate of the variance** $\mathbb{V}[\hat{\theta}_n]$ can be used for constructing a **confidence interval** for θ^* (not covered in the course).

Unbiasedness and Consistency

Note that

- the unbiasedness **does not** imply the consistency;
 - ▶ An **unbiased** estimator can be **inconsistent**.
- the consistency **does not** require the unbiasedness;
 - ▶ A **biased** estimator can be **consistent** (we'll see this later).

Example of an Unbiased Estimator that is not Consistent

Consider the **mean estimation** problem.

- Let $X \sim P$, and assume that $X_1, \dots, X_n \sim P$ (*i.i.d.*).
- Define an estimator F_n by

$$\hat{\mu} := F_n(X_1, \dots, X_n) := X_1.$$

i.e., we only use X_1 , and discard X_2, \dots, X_n .

- Then, this estimator is **unbiased**: In fact,

$$\mathbb{E}[\hat{\mu}] = \mathbb{E}[X_1] = \mathbb{E}[X] = \mu_P.$$

Example of an Unbiased Estimator that is not Consistent

- However, the variance of the estimate $\hat{\mu}$ is a constant:

$$\mathbb{V}[\hat{\mu}] = \mathbb{E}[(\hat{\mu} - \mathbb{E}[\hat{\mu}])^2] = \mathbb{E}[(X_1 - \mu_P)^2] = \mathbb{E}[(X - \mu_P)^2] = \sigma_P^2.$$

- Thus, the MSE of this estimator is

$$\mathbb{E}[(\hat{\mu} - \mu_P)^2] = \mathbb{E}[(\hat{\mu} - \mathbb{E}[\hat{\mu}])^2] = \sigma_P^2.$$

- Thus, the MSE does not decrease to 0, even if $n \rightarrow \infty$, i.e., the estimator is **not consistent**.

This example demonstrates that the **unbiasedness does not imply consistency**.

- For **consistency**, we need to make sure that the **variance of the estimate decreases to 0 as $n \rightarrow \infty$** .

Constructing an Unbiased Estimator by Weighting

Consider again the **mean estimation** problem.

- Let $X \sim P$, and assume $X_1, \dots, X_n \sim Q$ (*i.i.d.*).
- Assume that the data distribution Q is **different** from the target P .
- We show here that we can still construct an **unbiased** estimator of the mean

$$\mu_P = \mathbb{E}_{X \sim P}[X]$$

from the data $X_1, \dots, X_n \sim Q$ (*i.i.d.*).

Constructing an Unbiased Estimator by Weighting

- To this end, assume that distributions P and Q have density functions p and q , respectively.
- Define a weight function by

$$w(x) := \frac{p(x)}{q(x)}, \quad x \in \mathbb{R}$$

- Assume that this weight function is well-defined and bounded:

$$\max_{x \in \mathbb{R}} w(x) =: C < \infty.$$

- Note that this requires $p(x)/q(x) < C$, and thus

$$p(x) < Cq(x) \quad \text{for all } x \in \mathbb{R}.$$

- Thus, if the target density has a positive value $p(x) > 0$, then the data density should also have a positive value $q(x) > 0$.

Constructing an Unbiased Estimator by Weighting

- We assume for simplicity that this weight function $w(x) = p(x)/q(x)$ is **known**.

► Otherwise we need to estimate it from data.

- Define an estimator F_n of the mean μ_P as:

$$\hat{\mu} := F_n(X_1, \dots, X_n) := \frac{1}{n} \sum_{i=1}^n w(X_i) X_i.$$

- This is an **unbiased estimator** of the mean μ_P of P : This can be shown as follows.

Constructing an Unbiased Estimator by Weighting

- Recall that X_1, \dots, X_n are i.i.d. with $Y \sim Q$. Therefore,

$$\begin{aligned}\mathbb{E}[\hat{\mu}] &= \mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n w(X_i)X_i\right] \\&= \frac{1}{n} \sum_{i=1}^n \mathbb{E}[w(X_i)X_i] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}[w(Y)Y] \\&= \mathbb{E}[w(Y)Y] = \int x w(x) dQ(x) = \int x \frac{p(x)}{q(x)} q(x) dx \\&= \int x p(x) dx = \int x dP(x) = \mu_P.\end{aligned}$$

- Thus, $\hat{\mu} = \frac{1}{n} \sum_{i=1}^n w(X_i)X_i$ is an unbiased estimator of μ_P .

Constructing an Unbiased Estimator by Weighting

- On the other hand, the variance of the estimator is

$$\begin{aligned}\mathbb{V}[\hat{\mu}] &= \mathbb{V}\left[\frac{1}{n} \sum_{i=1}^n w(X_i)X_i\right] = \sum_{i=1}^n \mathbb{V}\left[\frac{1}{n} w(X_i)X_i\right] \\ &= \sum_{i=1}^n \frac{1}{n^2} \mathbb{V}[w(X_i)X_i] = \sum_{i=1}^n \frac{1}{n^2} \mathbb{V}[w(Y)Y] \\ &= \frac{1}{n} \mathbb{V}[w(Y)Y].\end{aligned}$$

- This can be upper-bounded as

$$\begin{aligned}\frac{1}{n} \mathbb{V}[w(Y)Y] &= \frac{1}{n} (\mathbb{E}[(w(Y)Y)^2] - (\mathbb{E}[w(Y)Y])^2) \\ &\leq \frac{1}{n} (\mathbb{E}[C^2 Y^2] + \mu_P^2) = \frac{1}{n} (C^2 \mathbb{E}[Y^2] - \mu_P^2) \\ &= \frac{1}{n} (C^2 (\sigma_Q^2 + \mu_Q^2) - \mu_P^2).\end{aligned}$$

Constructing an Unbiased Estimator by Weighting

- To summarize, the MSE of the estimator is

$$\mathbb{E}[(\hat{\mu} - \mu_P)^2] = \frac{1}{n} \mathbb{V}[w(Y)Y] \leq \frac{1}{n} (C^2(\sigma_Q^2 + \mu_Q^2) - \mu_P^2).$$

- Therefore, the MSE decreases to 0 as $n \rightarrow \infty$:

► i.e., the estimator $\hat{\mu} = \frac{1}{n} \sum_{i=1}^n w(X_i)X_i$ is **consistent** in estimating μ_P .

- The weight function $w(x)$ is called the **importance weight** of a point x .
- The way of constructing an estimator by weighting each sample point X_i by $w(X_i)$ is called **importance weighting**.

Constructing an Unbiased Estimator by Weighting

- Importance weighting is a widely used technique, examples including:
 - ▶ Domain shift adaptation in machine learning.
 - ▶ Estimation of treatment effects in causal inference.
 - ▶ Monte Carlo for efficient simulations.
- If you are interested in the first, you can for instance look at [Sugiyama and Kawanabe, 2012].

Outline

Mean Estimation Problem and Motivations

The Data Generation Process Matters

Preliminaries: Key Properties of Expectation and Variance

Statistical Estimators

Mean Square Error and Bias-Variance Decomposition

Bias-Variance Decomposition in Mean Estimation

Consistency and Unbiasedness

Variance Reduction by Introducing a Bias

Variance of Unbiased Estimators may be Large

- We demonstrate here that sometimes **biased estimators** may be “better” than **unbiased estimators**.
- The key is an approach called **shrinkage** or **regularization**, which is ubiquitous in statistics and machine learning.

Variance of Unbiased Estimators may be Large

- We have seen the bias-variance decomposition of the MSE:

$$\mathbb{E}[\|\hat{\theta}_n - \theta^*\|^2] = \underbrace{\mathbb{E}[\|\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n]\|^2]}_{\text{Variance}} + \underbrace{\|\mathbb{E}[\hat{\theta}_n] - \theta^*\|^2}_{\text{Bias}}$$

- The MSE decomposes into the **bias** and **variance**.
- For an **unbiased** estimator (i.e., the bias is zero), the MSE is equal to the variance:

$$\mathbb{E}[\|\hat{\theta}_n - \theta^*\|^2] = \underbrace{\mathbb{E}[\|\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n]\|^2]}_{\text{Variance}}.$$

- This variance may be **large** if, e.g.,
 - ▶ the sample size n is small
 - ▶ the dimensionality of $\hat{\theta}_n$ is large (in multivariate cases).
- In such a situation, a **biased** estimator with a **lower variance** may have a **smaller MSE** than the unbiased estimator.

Variance Reduction in Mean Estimation

- To describe this, consider the **mean estimation** problem.
- Let $X \sim P$, and $X_1, \dots, X_n \sim P$ (*i.i.d.*).
- We saw that the empirical average

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n X_i$$

is an **unbiased** estimator of the mean of

$$\mu_P := \mathbb{E}_{X \sim P}[X],$$

and the MSE is given by

$$\mathbb{E}[(\hat{\mu} - \mu_P)^2] = \frac{\mathbb{V}[X]}{n}.$$

- We'll show that there are **biased estimators** that have **smaller MSE** than the empirical average.

Empirical Average as a Least-Squares Solution

- We first show that the empirical average $\hat{\mu} = \frac{1}{n} \sum_{i=1}^n X_i$ is the solution to the following **optimization problem**

$$\hat{\mu} = \arg \min_{\alpha \in \mathbb{R}} \frac{1}{n} \sum_{i=1}^n (\alpha - X_i)^2.$$

- i.e., we consider a **least-squares problem** (fitting a constant α to the data X_1, \dots, X_n).
- To solve this, set the **derivative** of the objective function with respect to α to be zero:

$$\frac{d}{d\alpha} \left(\frac{1}{n} \sum_{i=1}^n (\alpha - X_i)^2 \right) = \frac{1}{n} \sum_{i=1}^n 2(\alpha - X_i) = 2\alpha - \frac{2}{n} \sum_{i=1}^n X_i = 0.$$

- Thus, the α that minimizes the objective function is

$$\alpha = \frac{1}{n} \sum_{i=1}^n X_i,$$

i.e., the empirical average.

Regularized Least Squares and Shrinkage Estimator

- We then consider a modified optimization problem, adding a **regularization term**:

$$\hat{\mu}_\lambda := \arg \min_{\alpha \in \mathbb{R}} \frac{1}{n} \sum_{i=1}^n (\alpha - X_i)^2 + \lambda \alpha^2,$$

where $\lambda \geq 0$ is a **regularization constant**.

- The solution is given by setting the **derivative** of the objective function to be 0:

$$\begin{aligned} \frac{d}{d\alpha} \left(\frac{1}{n} \sum_{i=1}^n (\alpha - X_i)^2 + \lambda \alpha^2 \right) &= \frac{1}{n} \sum_{i=1}^n 2(\alpha - X_i) + 2\lambda \alpha \\ &= 2\alpha - \frac{2}{n} \sum_{i=1}^n X_i + 2\lambda \alpha = 2\alpha(1 + \lambda) - \frac{2}{n} \sum_{i=1}^n X_i = 0. \end{aligned}$$

- Thus, the solution is given by

$$\alpha = \frac{1}{(1 + \lambda)} \frac{1}{n} \sum_{i=1}^n X_i =: \hat{\mu}_\lambda$$

Regularized Least Squares and Shrinkage Estimator

$$\hat{\mu}_{\lambda} = \frac{1}{(1 + \lambda)} \frac{1}{n} \sum_{i=1}^n X_i.$$

- Large λ **shrinks** the solution $\hat{\mu}_{\lambda}$ towards 0.
 - ▶ In this sense, this is called a **shrinkage estimator**.
- $\lambda = 0$ recovers the empirical average $\hat{\mu}_0 = \frac{1}{n} \sum_{i=1}^n X_i$.

Mean Square Error of the Shrinkage Estimator

- The expectation of $\hat{\mu}_\lambda$ is

$$\mathbb{E}[\hat{\mu}_\lambda] = \mathbb{E}\left[\frac{1}{(1+\lambda)} \frac{1}{n} \sum_{i=1}^n X_i\right] = \frac{1}{(1+\lambda)} \mu_P.$$

- Thus, the (squared) **bias** of $\hat{\mu}_\lambda$ is

$$(\mathbb{E}[\hat{\mu}_\lambda] - \mu_P)^2 = \left(\frac{1}{(1+\lambda)} \mu_P - \mu_P\right)^2 = \frac{\lambda^2 \mu_P^2}{(1+\lambda)^2}.$$

- Thus, the **bias increases** as λ **increases**.

- On the other hand, the **variance** of $\hat{\mu}_\lambda$ is

$$\begin{aligned} \mathbb{V}[\hat{\mu}_\lambda] &= \mathbb{V}\left[\frac{1}{1+\lambda} \frac{1}{n} \sum_{i=1}^n X_i\right] \\ &= \frac{1}{(1+\lambda)^2} \mathbb{V}\left[\frac{1}{n} \sum_{i=1}^n X_i\right] = \frac{1}{(1+\lambda)^2} \frac{\mathbb{V}[X]}{n}. \end{aligned}$$

- Thus, the **variance decreases** as λ **increases**.

Mean Square Error of the Shrinkage Estimator

- Thus, the MSE of $\hat{\mu}_\lambda$ is

$$\begin{aligned}\mathbb{E}[(\hat{\mu}_\lambda - \mu_P)^2] &= \mathbb{V}[\hat{\mu}_\lambda] + (\mathbb{E}[\hat{\mu}_\lambda] - \mu_P)^2 \\ &= \frac{1}{(1 + \lambda)^2} \frac{\mathbb{V}[X]}{n} + \frac{\lambda^2 \mu_P^2}{(1 + \lambda)^2}\end{aligned}$$

- Let's draw some observations. Assume $\mu_P \neq 0$.

Mean Square Error of the Shrinkage Estimator

- By an easy calculation, the MSE of $\hat{\mu}_\lambda = \frac{1}{(1+\lambda)} \frac{1}{n} \sum_{i=1}^n X_i$ can be shown to be **smaller** than that of the empirical average

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n X_i:$$

$$\mathbb{E}[(\hat{\mu}_\lambda - \mu_P)^2] < \mathbb{E}[(\hat{\mu} - \mu_P)^2]$$

if $\lambda > 0$ is chosen so that

$$\frac{\lambda}{2 + \lambda} \leq \frac{\mathbb{V}[X]}{n \mu_P^2}.$$

Some interpretations:

- When $\mathbb{V}[X]/n$ is large (e.g., when n is small), a large λ can be taken (and more shrinkage).
- When the mean μ_P^2 is small, a large λ can be taken (and more shrinkage).

Exercise: Perform numerical experiments to confirm that the shrinkage estimator can have a smaller MSE.

Mean Square Error of the Shrinkage Estimator

- For a right choice of $\lambda > 0$, we need to know $\mathbb{V}[X]$ and μ_P .
 - ▶ Therefore this estimator is not practically useful.
- However, under some assumptions (e.g., P is a Gaussian), there is a way of choosing λ **without the knowledge** of $\mathbb{V}[X]$ and μ_P .
 - ▶ This resulting estimator is called the **James-Stein estimator**; see [Efron and Hastie, 2016, Section 7] [Berger, 1985, Section 5.4].

Regularization for Variance Reduction

- Anyway, this example illustrates that **artificially introducing a bias** is often useful to **reduce the variance**.
- In this spirit, **regularization** has been widely used in many statistical methods: e.g.,
 - ▶ L_2 and L_1 regularization in regression and classification (supervised learning)
 - ▶ Early stopping in optimization algorithms for machine learning algorithms.
- In supervised learning problems, a good regularization constant can be chosen by, e.g., **cross validation**
 - ▶ See e.g. the MALIS and ASI courses.

Summary of the Lecture

- We introduced several important concepts in statistical estimation.
- When constructing statistical estimators, always pay attention to
 - ▶ what is **your quantity of interest** (in the **population**).
 - ▶ **how your data were generated**.
 - ▶ whether your estimator is **biased** or **unbiased**.
 - ▶ how much your estimate would have **variance**.



Berger, J. O. (1985).

Statistical Decision Theory and Bayesian Analysis.

Springer Science & Business Media.



Efron, B. and Hastie, T. (2016).

Computer Age Statistical Inference.

Cambridge University Press.



Sugiyama, M. and Kawanabe, M. (2012).

Machine Learning in Non-Stationary Environments: Introduction to Covariate Shift Adaptation.

MIT Press, Cambridge, MA, USA.